


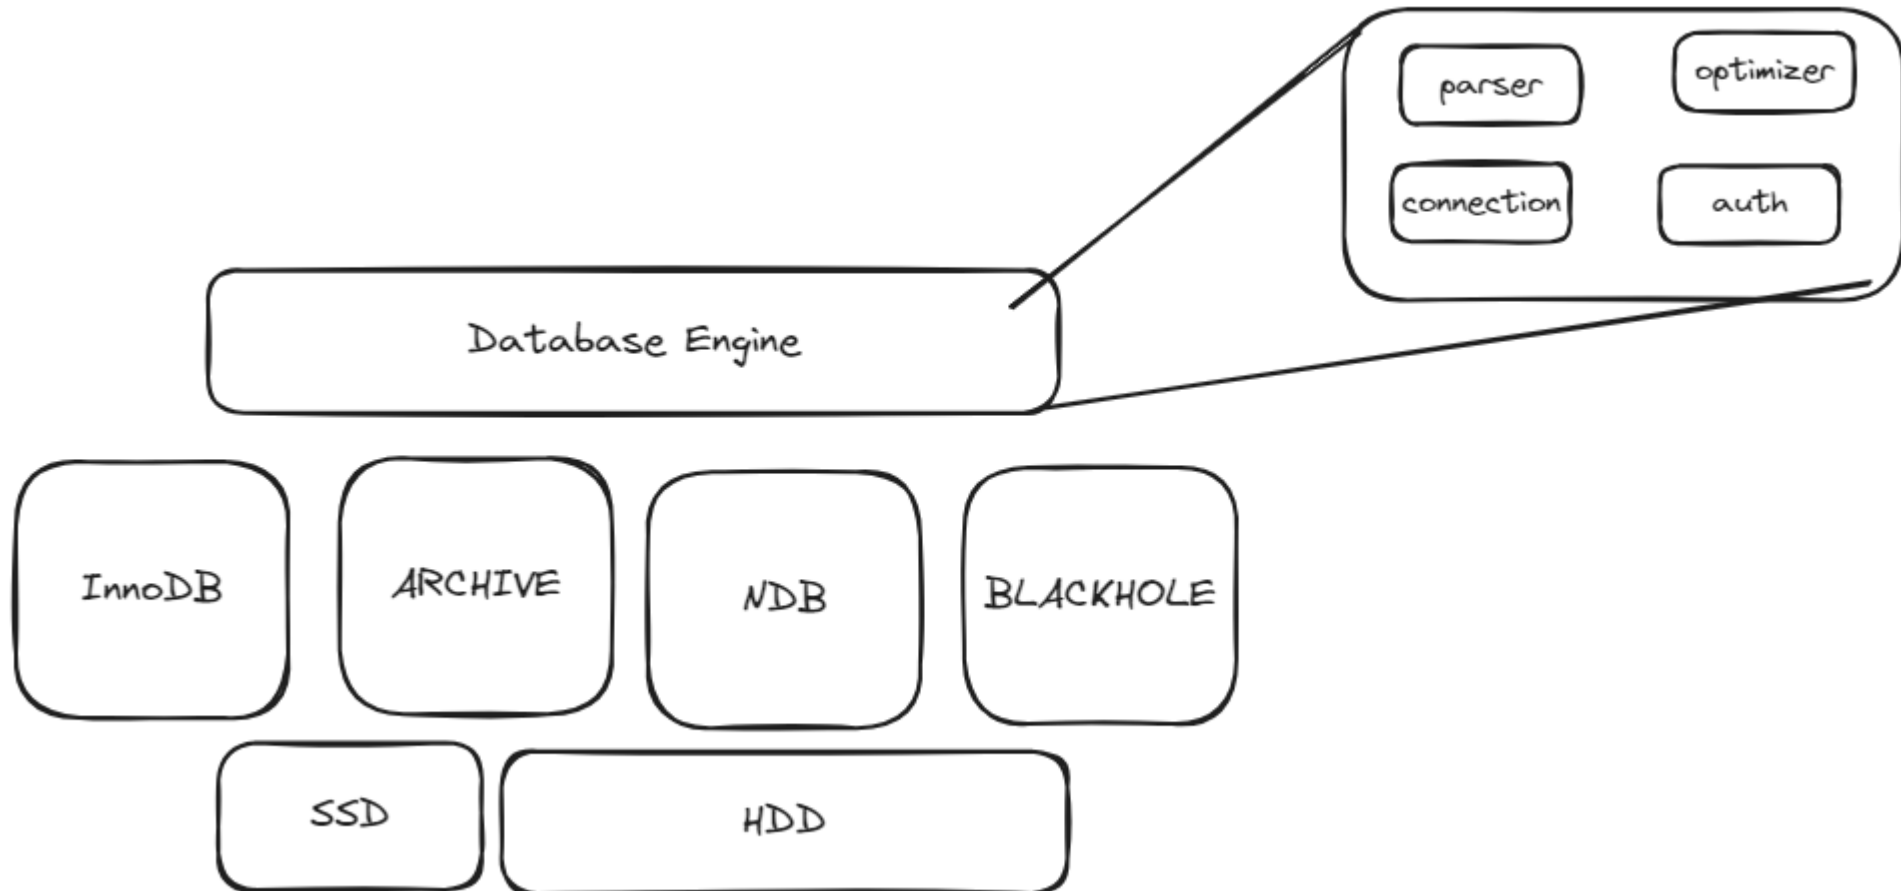
# Database internals

MVCC


# Database Engine (MySQL InnoDB)



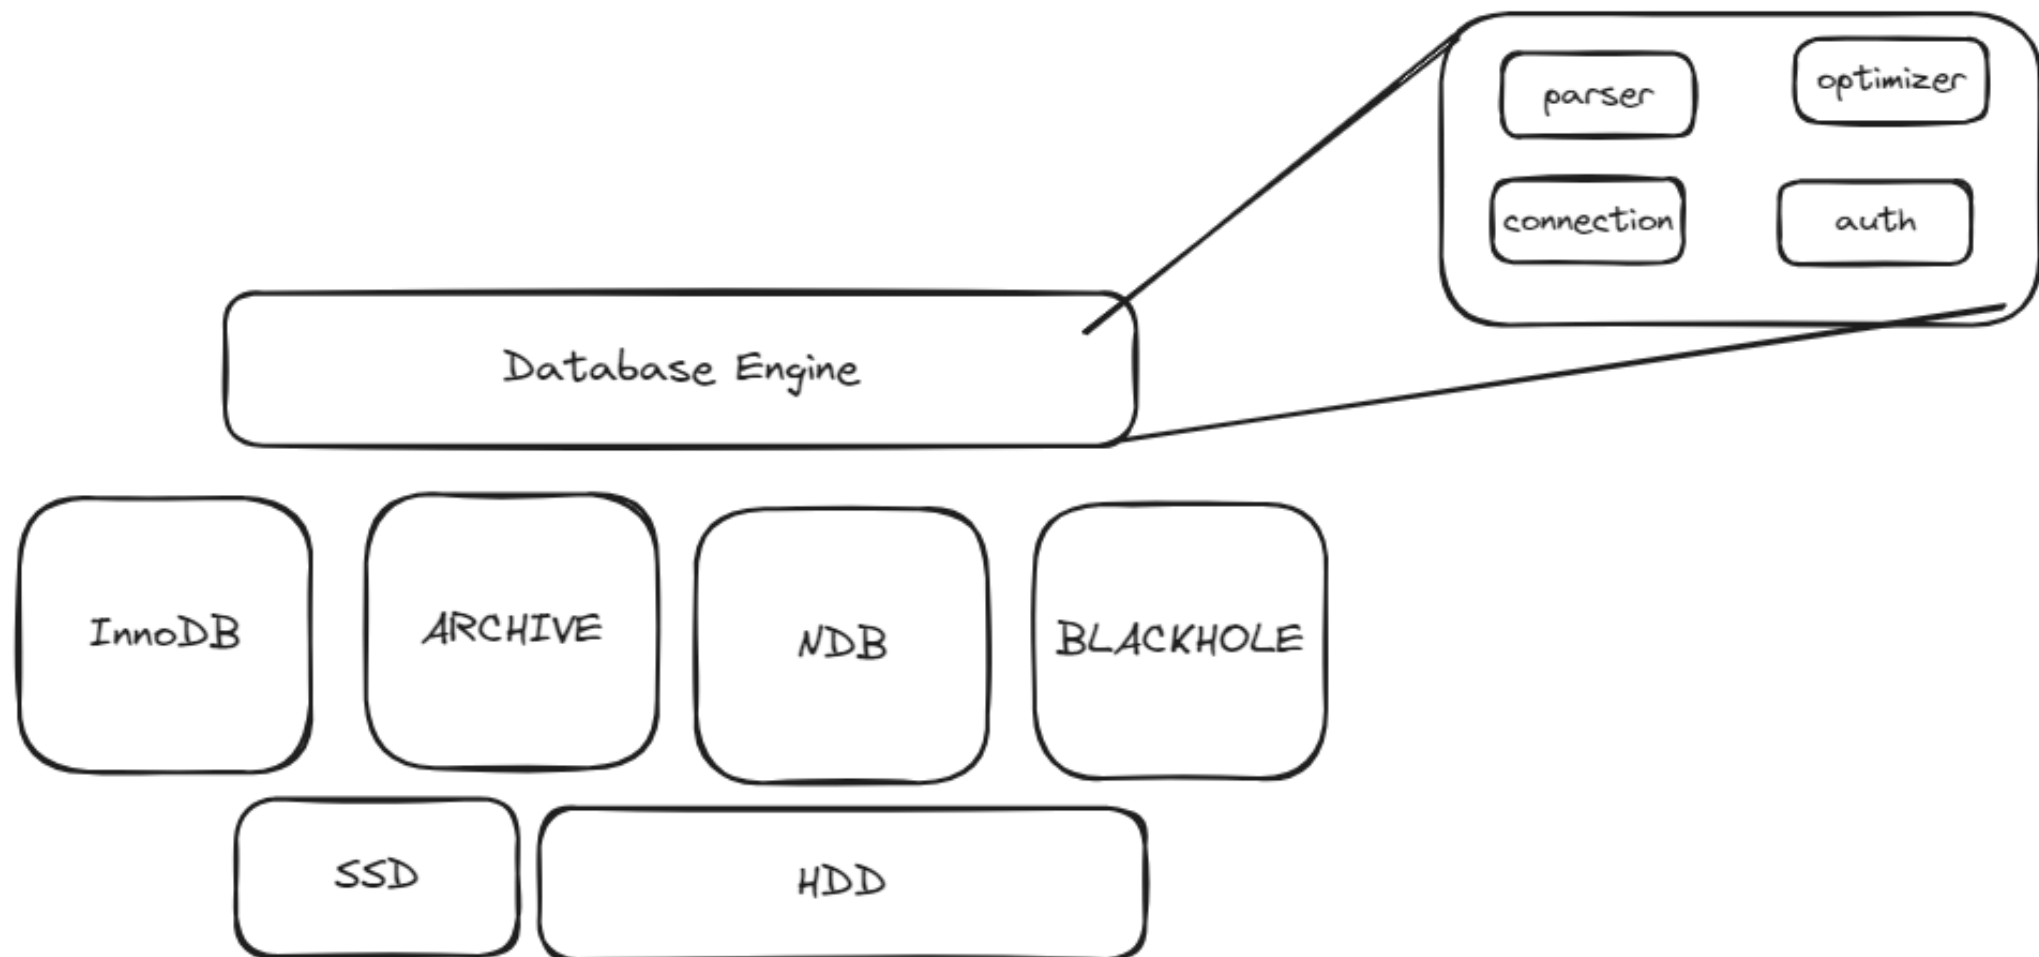
# Why should we use multi storage engine?



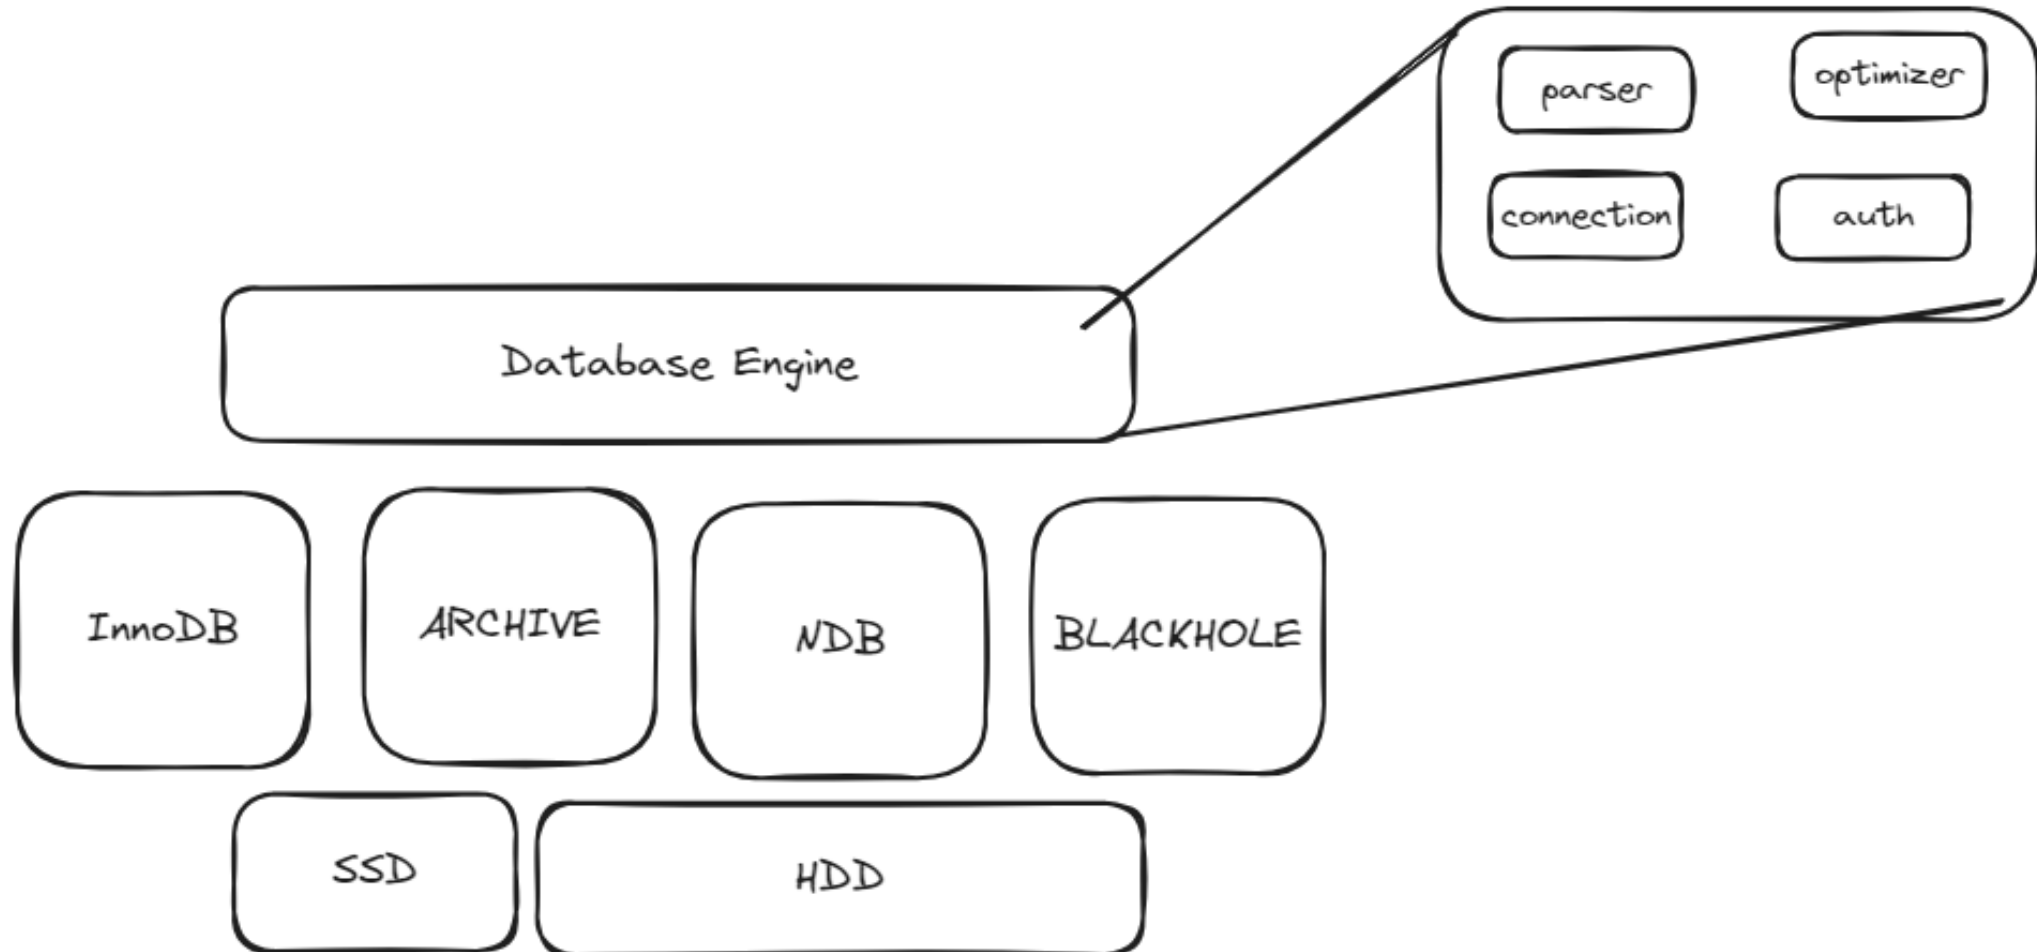
# Does storage engine share storage?



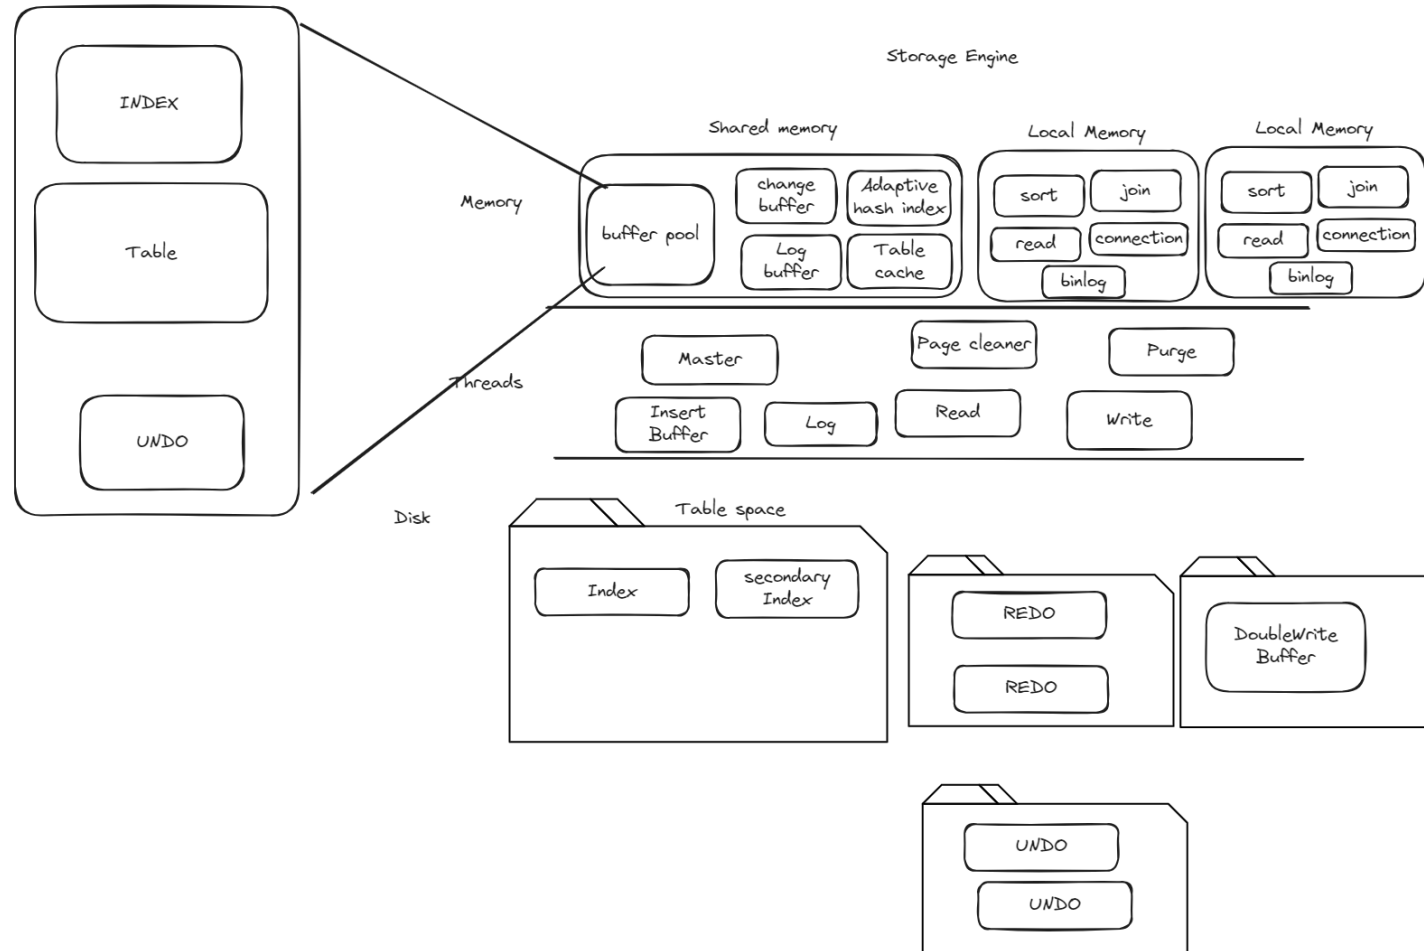
# OLTP vs OLAP



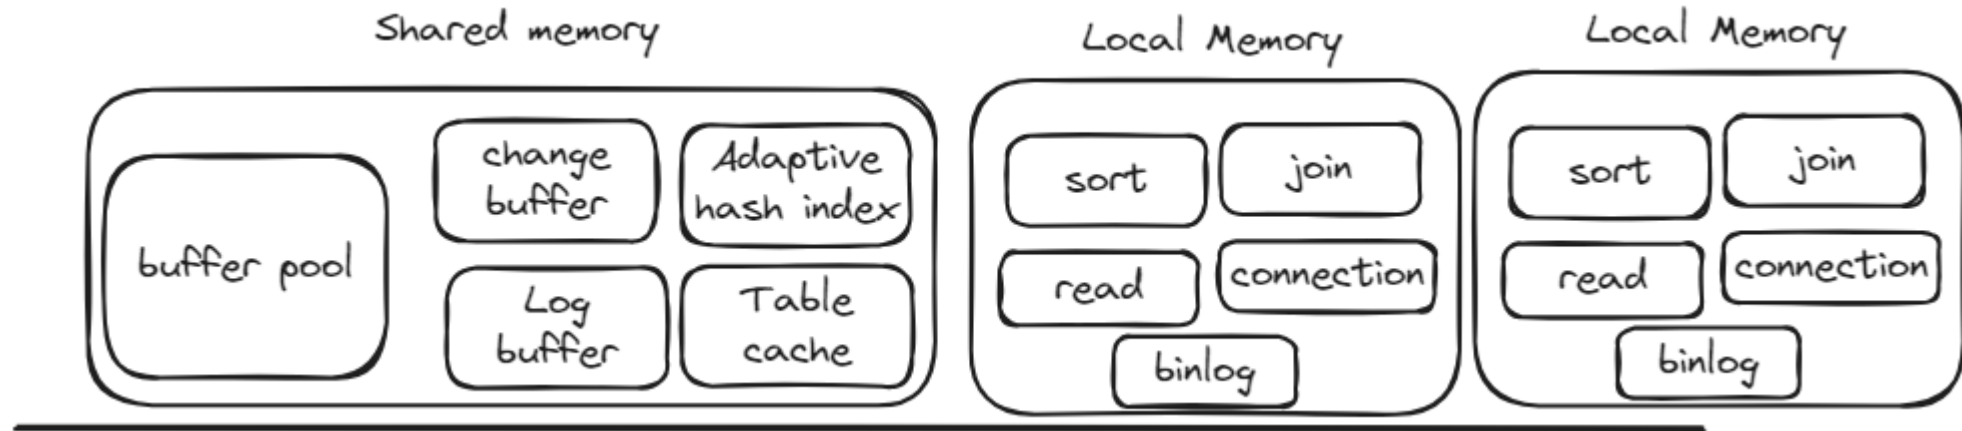
# How about using memory engine instead of Redis?



# How storage engine works?

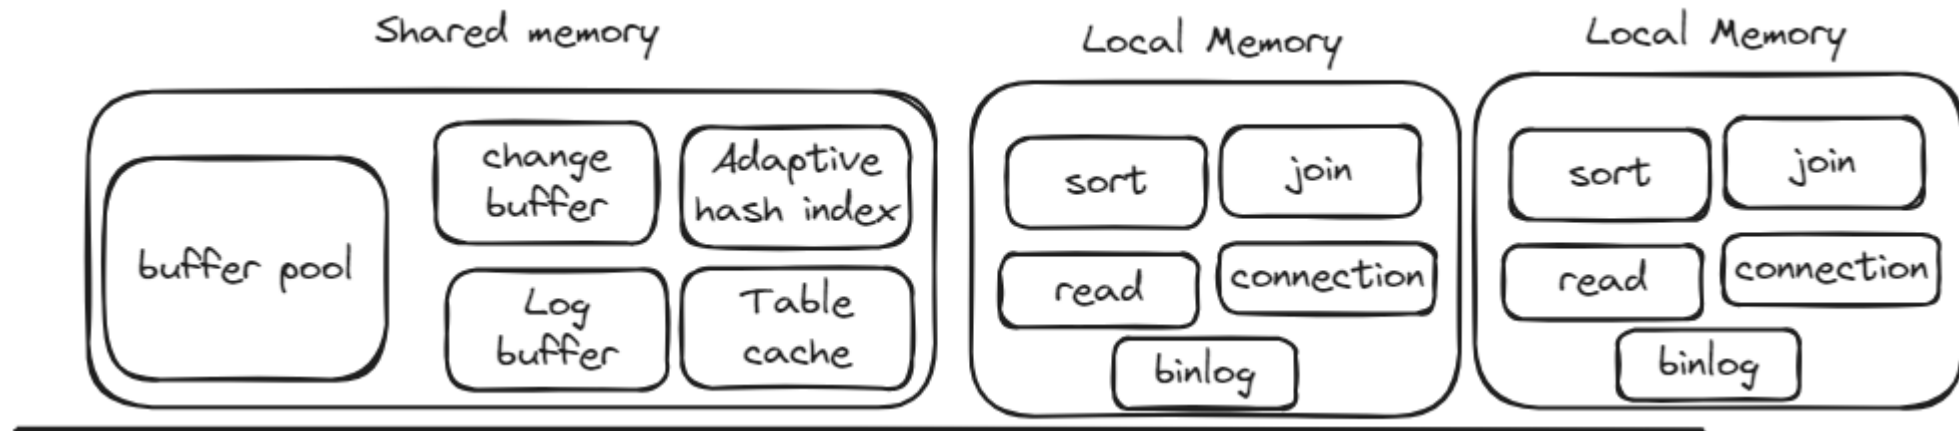


# Memory

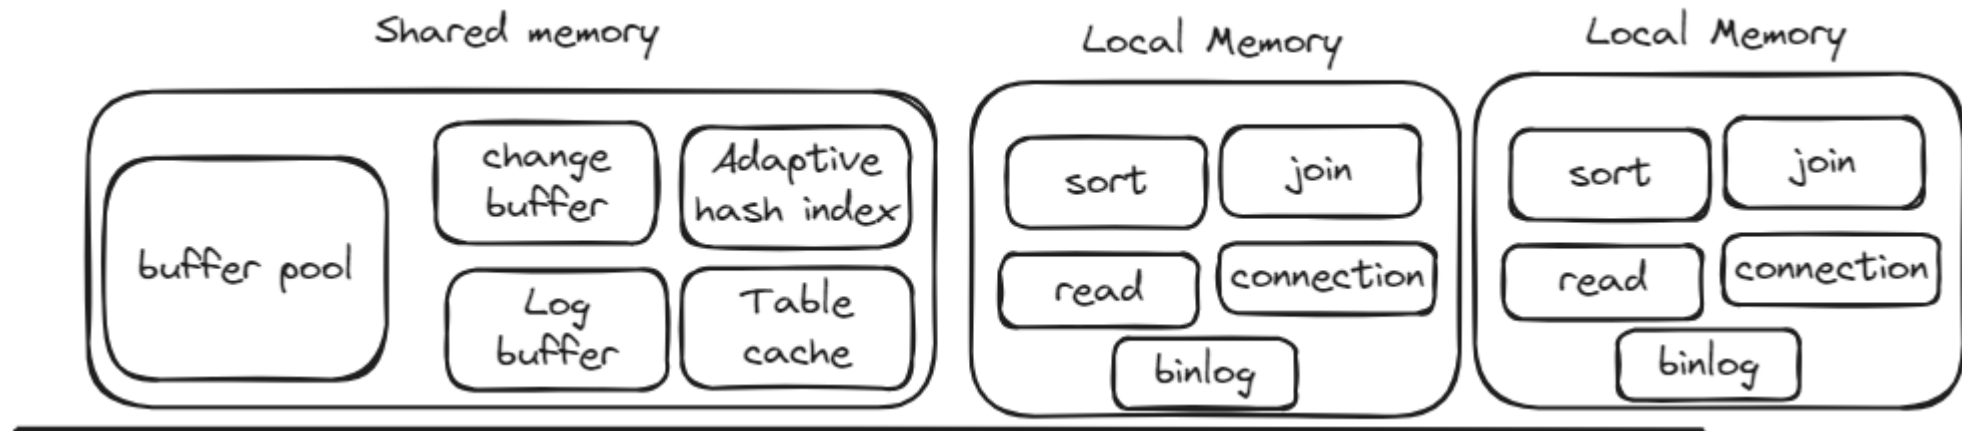





# Database connection?




# What's doing DBA?




# Buffer pool



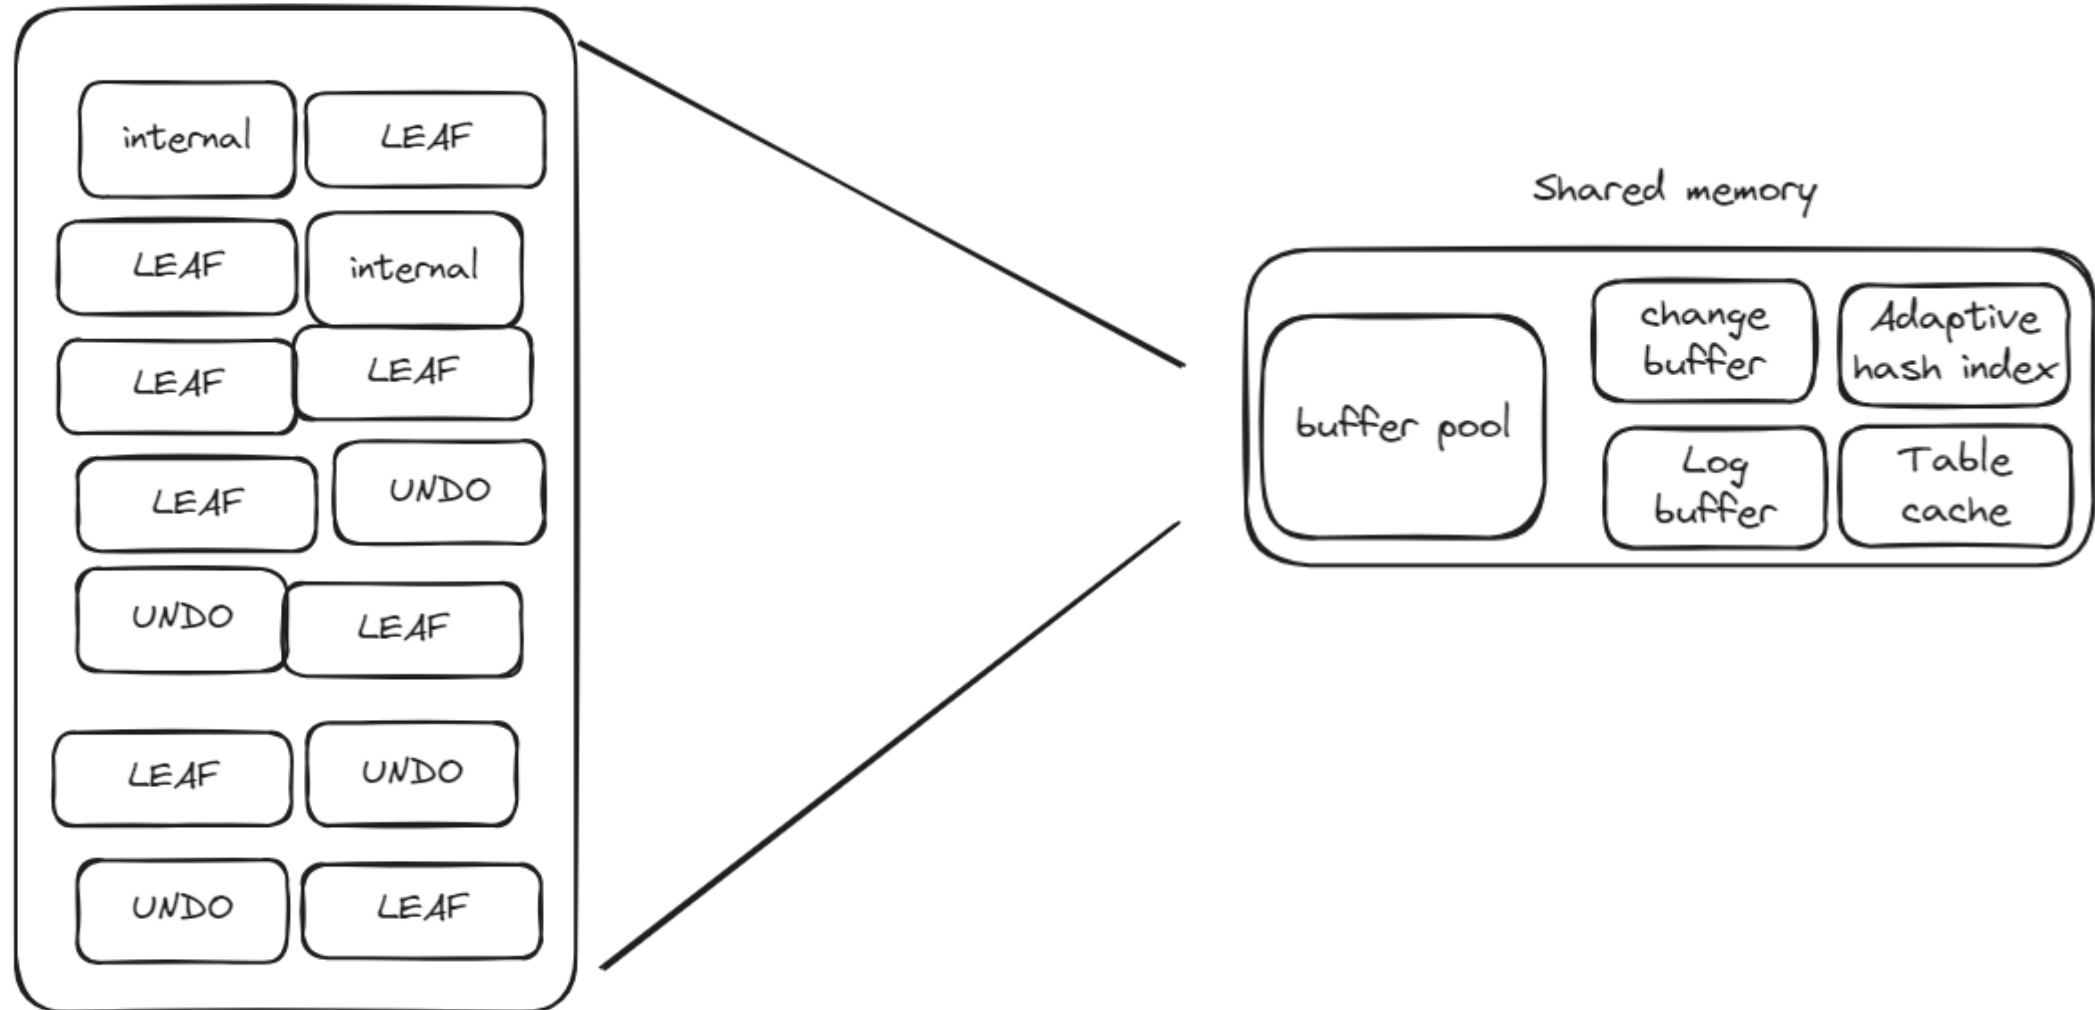
# If Commit, data is written on disk?




# What is dirty page??



# How data cleaned?



# What is Undo log?



# Isolation Level

**1.READ UNCOMMITTED**


**2.READ COMMITTED**

**3.REPEATABLE READ** (Default)


**4.SERIALIZABLE**




# Read uncommitted



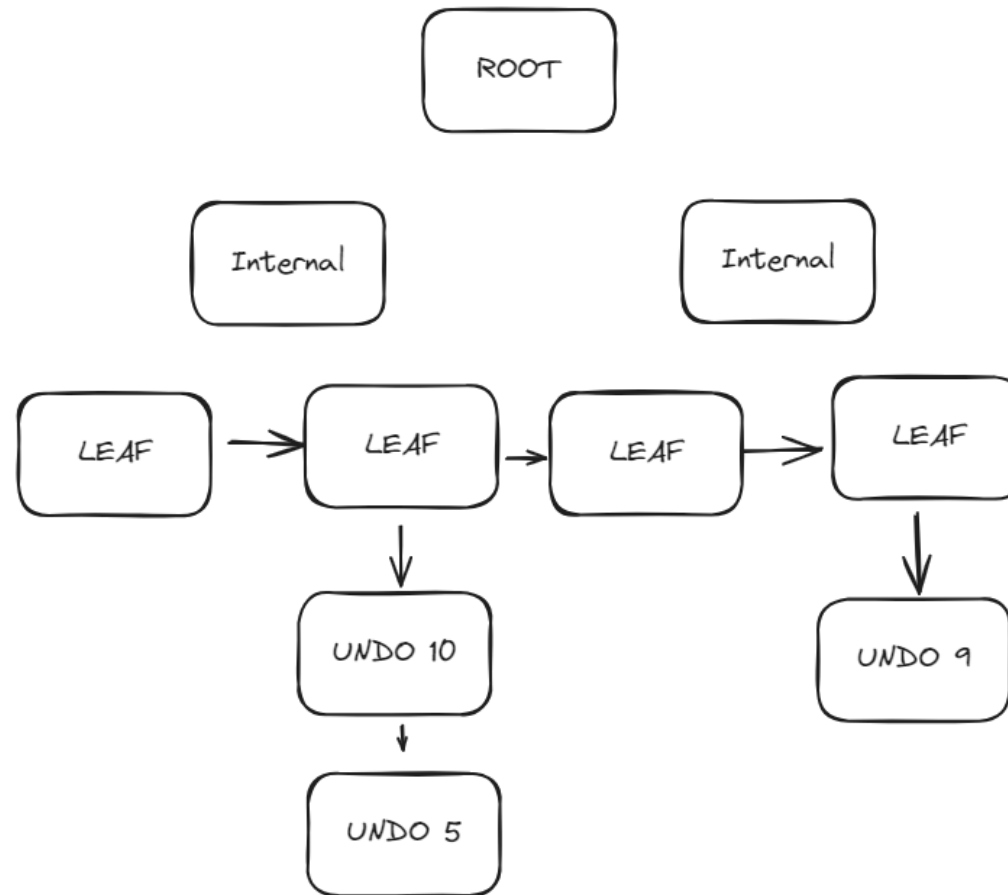
# What is recorded in leaf node?




# Leaf node and undo log



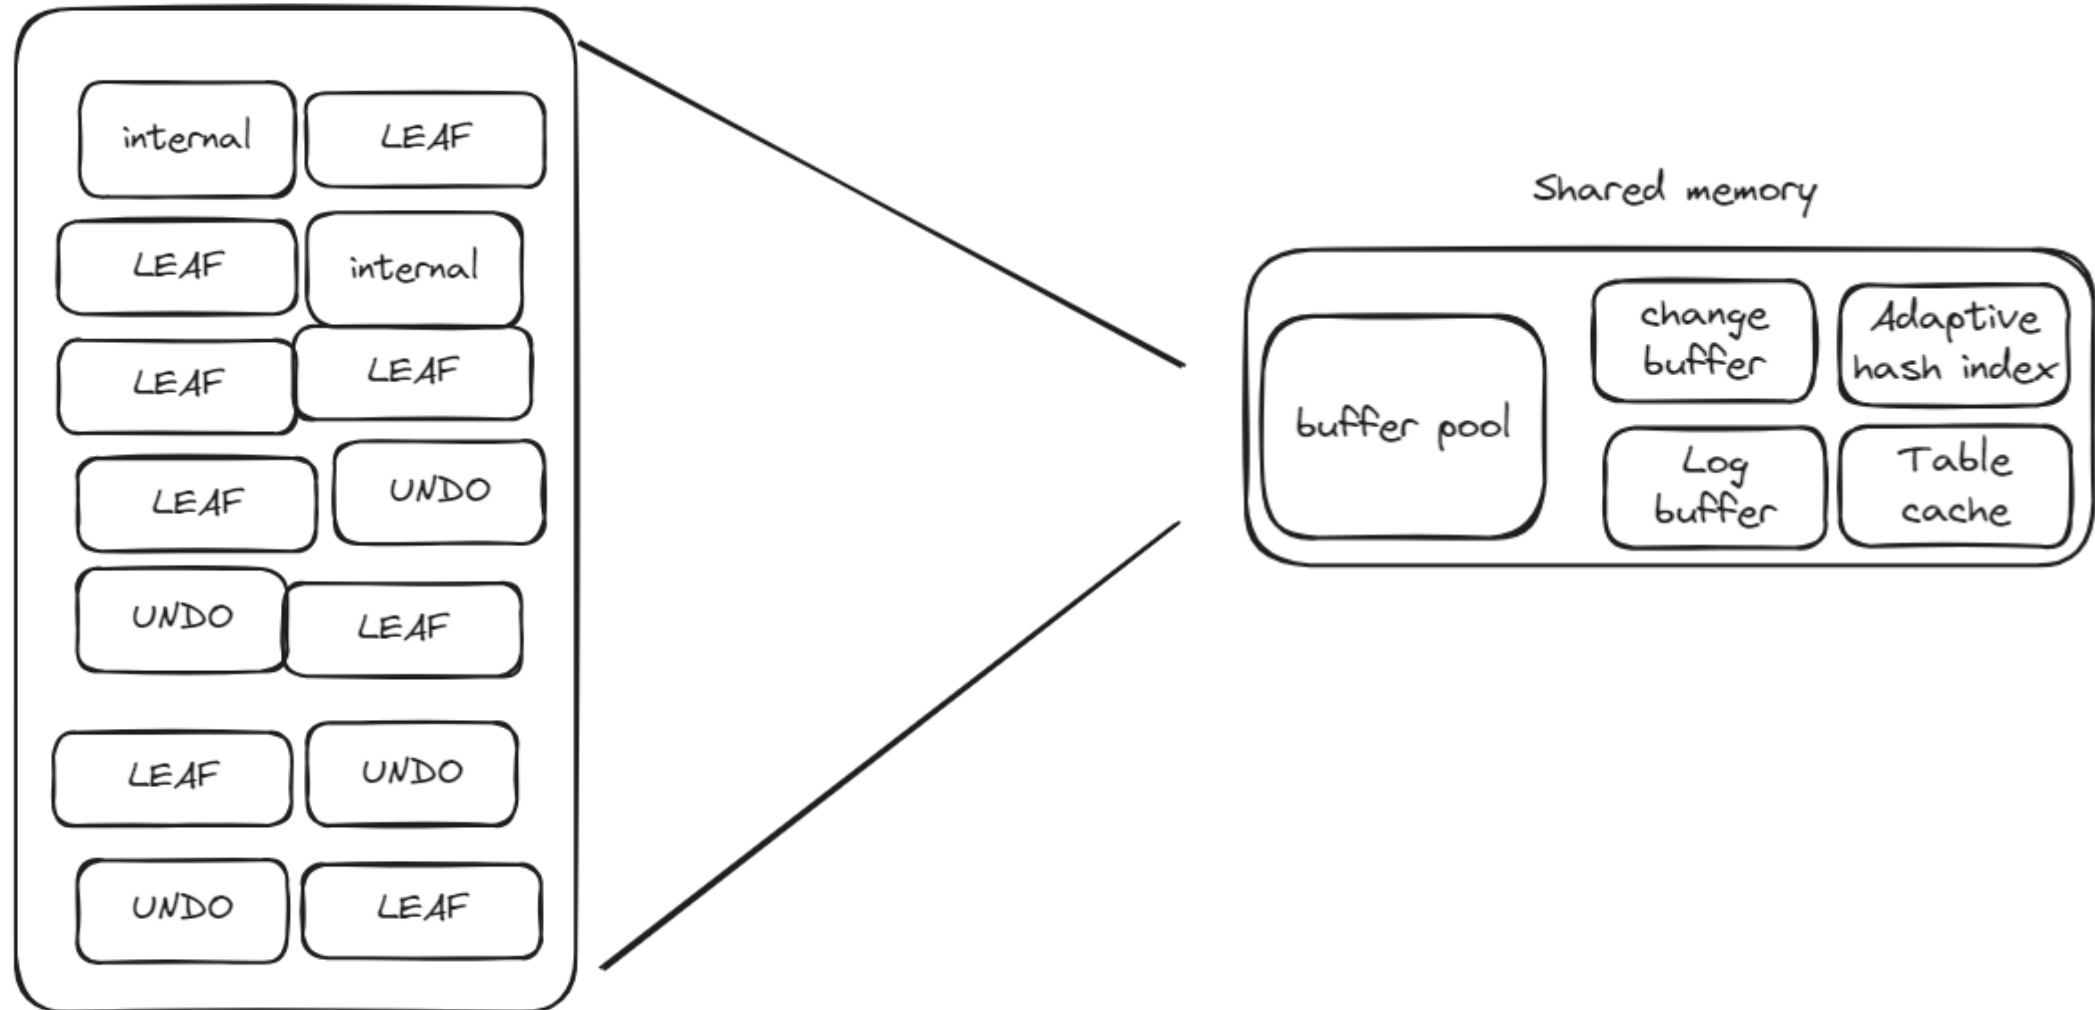
# What is undo log?




# Could read uncommitted use rollback?



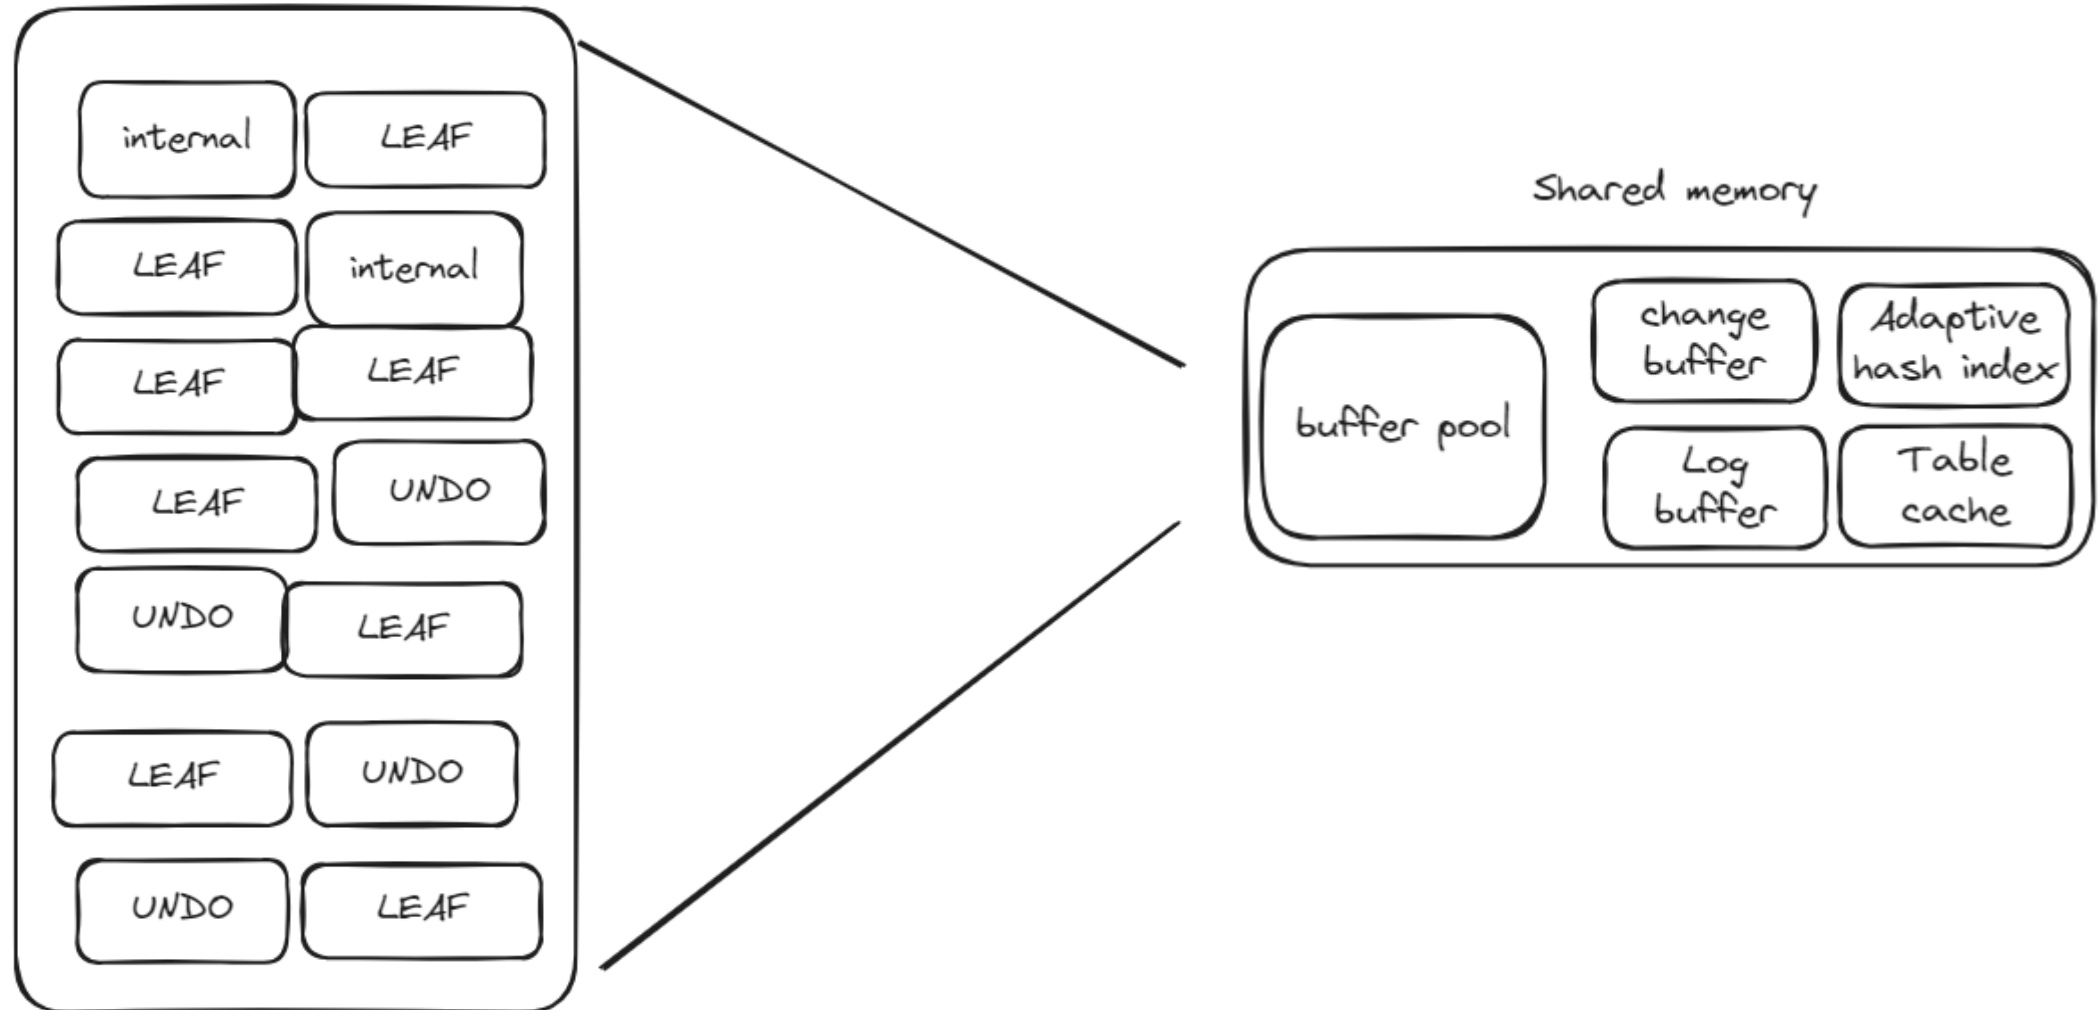
# Read committed



# Leaf node could be uncommitted

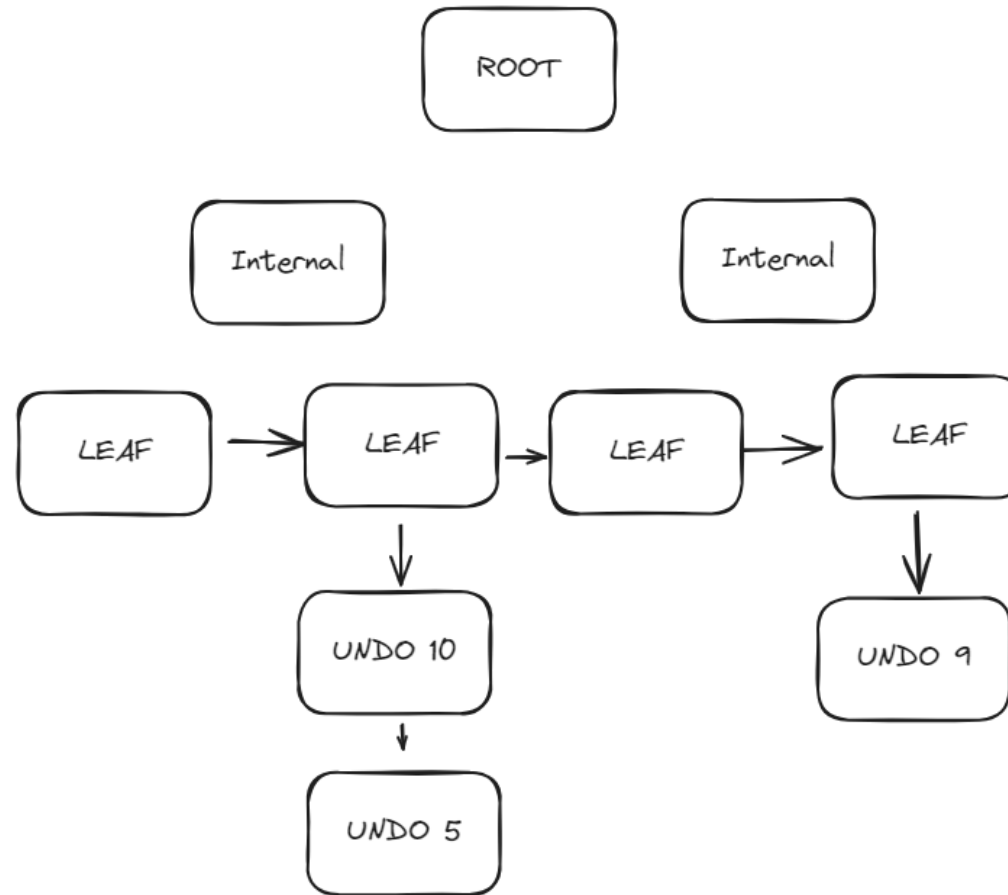


# Repeatable read







# Next key lock?



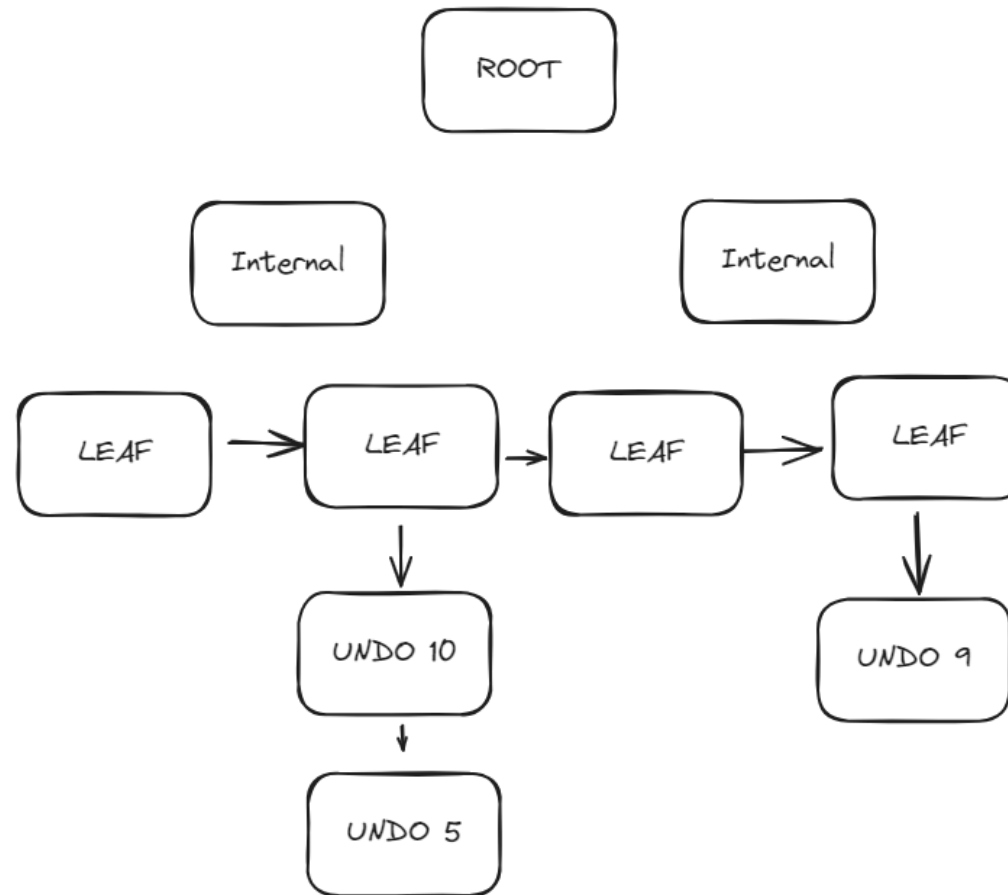
# Why deadlock occurs?



# serializable




# What's difference from repeatable read

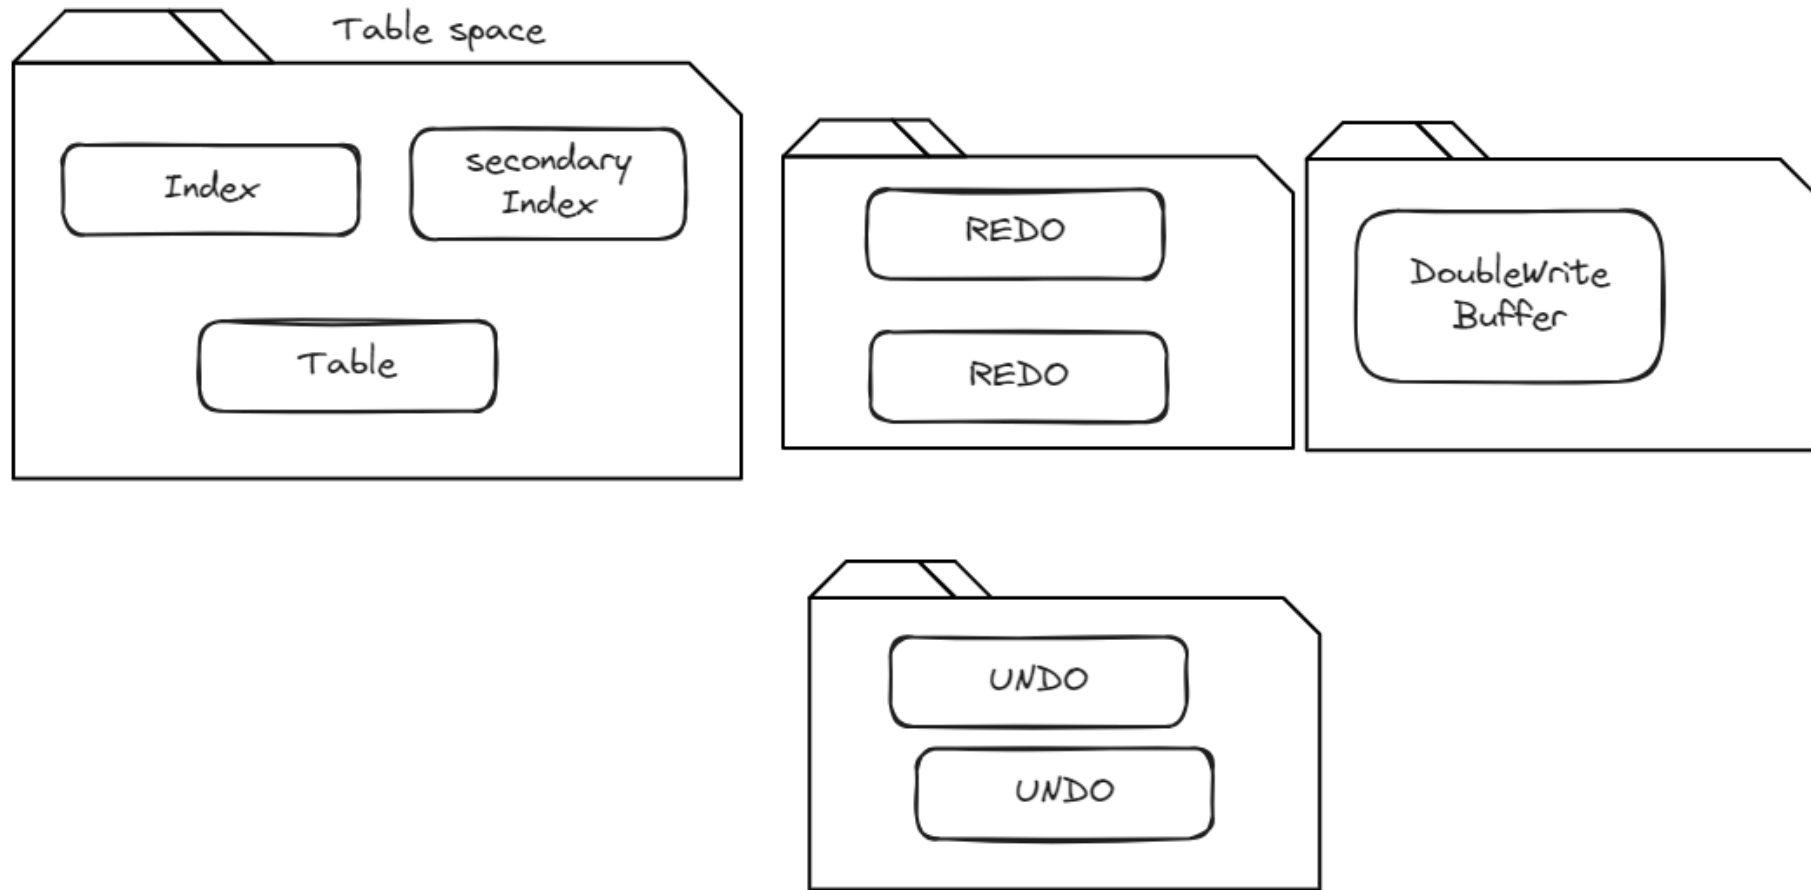


# Threads



Threads



# Storage



# Sharding vs replication



# How to make sharding scalable?

